
This is a companion notebook for the book [Deep Learning with Python, Second Edition](#). For readability, it only contains runnable code blocks and section titles, and omits everything else in the book: text paragraphs, figures, and pseudocode.

If you want to be able to follow what's going on, I recommend reading the notebook side by side with your copy of the book.

This notebook was generated for TensorFlow 2.6.

▼ Getting started with neural networks: Classification and regression

▼ Classifying movie reviews: A binary classification example

▼ The IMDB dataset

Loading the IMDB dataset

▼ Importing an IMDB dataset from Keras. Here, we'll look at the 10000 words.

▼ Dividing the dataset into training and test sets.

```
from tensorflow.keras.datasets import imdb
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(
    num_words=10000)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz
17464789/17464789 [=====] - 0s 0us/step
```

▼ Simply printing the first review from the training dataset.

```
train_data[0]
```



```

92,
25,
104,
4,
226,
65,
16,
38,
1334,
88,
12,
16,
283,
5,
16,
4472,
113,
103,
32,
15,
16,
5345,
19,
178,
321

```

▼ checking the first review's label

```
train_labels[0]
```

```
1
```

```
max([max(sequence) for sequence in train_data])
```

```
9999
```

Decoding and displaying movie reviews in text

```

word_index = imdb.get_word_index()
reverse_word_index = dict(
    [(value, key) for (key, value) in word_index.items()])
decoded_review = " ".join(
    [reverse_word_index.get(i - 3, "?") for i in train_data[0]])
#As can be seen, the first review is positive, and the label is 1.

```

```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb\_word\_index.json
1641221/1641221 [=====] - 0s 0us/step

```

▼ Preparing the data

Encoding the integer sequences via multi-hot encoding

```

import numpy as np
def vectorize_sequences(sequences, dimension=10000):
    results = np.zeros((len(sequences), dimension))
    for i, sequence in enumerate(sequences):
        for j in sequence:
            results[i, j] = 1.
    return results
x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)

```

```
#
```

```
x_train[0]
```

```
array([0., 1., 1., ..., 0., 0., 0.])
```

```

y_train = np.asarray(train_labels).astype("float32")
y_test = np.asarray(test_labels).astype("float32")

```

▼ Building your model

Model definition

```
from tensorflow import keras
from tensorflow.keras import layers
# #Here I am using two hidden layers, each with 16 nodes, and only one node in the output layer for either +ve or -ve output. ReLu is used f
model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
```

Compiling the model

▼ Adam is used as the optimizer, and binary crossentropy is used as the loss function.

```
model.compile(optimizer="adam",
              loss="binary_crossentropy",
              metrics=["accuracy"])
```

▼ Validating your approach

Setting aside a validation set

```
x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

Training your model

▼ we're training our model with 20 epochs and 512 batches.

```
history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=20,
                    batch_size=512,
                    validation_data=(x_val, y_val))
```

```
Epoch 1/20
30/30 [=====] - 4s 90ms/step - loss: 0.5607 - accuracy: 0.7707 - val_loss: 0.4155 - val_accuracy: 0.8602
Epoch 2/20
30/30 [=====] - 1s 41ms/step - loss: 0.3134 - accuracy: 0.9019 - val_loss: 0.3038 - val_accuracy: 0.8886
Epoch 3/20
30/30 [=====] - 1s 49ms/step - loss: 0.2096 - accuracy: 0.9337 - val_loss: 0.2811 - val_accuracy: 0.8878
Epoch 4/20
30/30 [=====] - 2s 58ms/step - loss: 0.1543 - accuracy: 0.9533 - val_loss: 0.2805 - val_accuracy: 0.8883
Epoch 5/20
30/30 [=====] - 1s 40ms/step - loss: 0.1165 - accuracy: 0.9685 - val_loss: 0.2963 - val_accuracy: 0.8834
Epoch 6/20
30/30 [=====] - 1s 33ms/step - loss: 0.0893 - accuracy: 0.9771 - val_loss: 0.3157 - val_accuracy: 0.8820
Epoch 7/20
30/30 [=====] - 2s 51ms/step - loss: 0.0683 - accuracy: 0.9857 - val_loss: 0.3422 - val_accuracy: 0.8794
Epoch 8/20
30/30 [=====] - 1s 42ms/step - loss: 0.0524 - accuracy: 0.9915 - val_loss: 0.3694 - val_accuracy: 0.8775
Epoch 9/20
30/30 [=====] - 1s 35ms/step - loss: 0.0396 - accuracy: 0.9939 - val_loss: 0.4013 - val_accuracy: 0.8758
Epoch 10/20
30/30 [=====] - 1s 34ms/step - loss: 0.0294 - accuracy: 0.9967 - val_loss: 0.4311 - val_accuracy: 0.8758
Epoch 11/20
30/30 [=====] - 1s 33ms/step - loss: 0.0224 - accuracy: 0.9987 - val_loss: 0.4591 - val_accuracy: 0.8733
Epoch 12/20
30/30 [=====] - 1s 32ms/step - loss: 0.0171 - accuracy: 0.9992 - val_loss: 0.4898 - val_accuracy: 0.8730
Epoch 13/20
30/30 [=====] - 1s 36ms/step - loss: 0.0134 - accuracy: 0.9995 - val_loss: 0.5123 - val_accuracy: 0.8714
```

```

Epoch 14/20
30/30 [=====] - 1s 49ms/step - loss: 0.0104 - accuracy: 0.9998 - val_loss: 0.5384 - val_accuracy: 0.8705
Epoch 15/20
30/30 [=====] - 2s 58ms/step - loss: 0.0082 - accuracy: 0.9999 - val_loss: 0.5612 - val_accuracy: 0.8697
Epoch 16/20
30/30 [=====] - 1s 38ms/step - loss: 0.0067 - accuracy: 0.9999 - val_loss: 0.5832 - val_accuracy: 0.8692
Epoch 17/20
30/30 [=====] - 1s 41ms/step - loss: 0.0055 - accuracy: 0.9999 - val_loss: 0.6048 - val_accuracy: 0.8688
Epoch 18/20
30/30 [=====] - 1s 34ms/step - loss: 0.0047 - accuracy: 0.9999 - val_loss: 0.6208 - val_accuracy: 0.8668
Epoch 19/20
30/30 [=====] - 1s 35ms/step - loss: 0.0040 - accuracy: 1.0000 - val_loss: 0.6398 - val_accuracy: 0.8666
Epoch 20/20
30/30 [=====] - 1s 32ms/step - loss: 0.0035 - accuracy: 1.0000 - val_loss: 0.6563 - val_accuracy: 0.8663

```

```

history_dict = history.history
history_dict.keys()

dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])


```

Plotting the training and validation loss

```

import matplotlib.pyplot as plt
history_dict = history.history
loss_values = history_dict["loss"]
val_loss_values = history_dict["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, "bo", label="Training loss")
plt.plot(epochs, val_loss_values, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()

```




Plotting the training and validation accuracy

```

plt.clf()
acc = history_dict["accuracy"]
val_acc = history_dict["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()

```



Validation accuracy starts to decline around the third epoch.

Retraining a model from scratch

```
model = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
#Here i am using three epochs to retrain the model here.
model.compile(optimizer="adam",
              loss="binary_crossentropy",
              metrics=["accuracy"])
model.fit(x_train, y_train, epochs=4, batch_size=512)
results = model.evaluate(x_test, y_test)
```

```
Epoch 1/4
49/49 [=====] - 2s 33ms/step - loss: 0.4773 - accuracy: 0.8111
Epoch 2/4
49/49 [=====] - 2s 36ms/step - loss: 0.2437 - accuracy: 0.9132
Epoch 3/4
49/49 [=====] - 1s 27ms/step - loss: 0.1790 - accuracy: 0.9363
Epoch 4/4
49/49 [=====] - 1s 28ms/step - loss: 0.1421 - accuracy: 0.9517
782/782 [=====] - 2s 2ms/step - loss: 0.3079 - accuracy: 0.8802
```

```
results

[0.3079053461551666, 0.8801599740982056]
```

▼ Building your model

▼ 1 using one or three hidden layers, and see how doing so

affects validation and test accuracy.

```
#I am creating a model with just 1 hidden layer and the ReLu activation function.
model1_1 = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])

# I am using three hidden layers here, with ReLu activation function and sigmoid for output layer.
model1_3 = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
```

```

layers.Dense(1, activation="sigmoid")
])

```

#Adam and binary crossentropy are used in both scenarios (3 and 1 layers)

```

model1_1.compile(optimizer="adam",
                 loss="binary_crossentropy",
                 metrics=["accuracy"])

```

```

model1_3.compile(optimizer="adam",
                 loss="binary_crossentropy",
                 metrics=["accuracy"])

```

model fitting with 20 epochs and 512 batch size

```

history1_1 = model1_1.fit(partial_x_train,
                          partial_y_train,
                          epochs=20,
                          batch_size=512,
                          validation_data=(x_val, y_val))

```

```

Epoch 1/20
30/30 [=====] - 4s 109ms/step - loss: 0.5741 - accuracy: 0.7357 - val_loss: 0.4392 - val_accuracy: 0.8503
Epoch 2/20
30/30 [=====] - 1s 43ms/step - loss: 0.3491 - accuracy: 0.8931 - val_loss: 0.3345 - val_accuracy: 0.8821
Epoch 3/20
30/30 [=====] - 1s 36ms/step - loss: 0.2585 - accuracy: 0.9230 - val_loss: 0.3004 - val_accuracy: 0.8854
Epoch 4/20
30/30 [=====] - 1s 33ms/step - loss: 0.2099 - accuracy: 0.9386 - val_loss: 0.2833 - val_accuracy: 0.8896
Epoch 5/20
30/30 [=====] - 1s 38ms/step - loss: 0.1765 - accuracy: 0.9501 - val_loss: 0.2800 - val_accuracy: 0.8857
Epoch 6/20
30/30 [=====] - 1s 35ms/step - loss: 0.1518 - accuracy: 0.9581 - val_loss: 0.2789 - val_accuracy: 0.8872
Epoch 7/20
30/30 [=====] - 2s 52ms/step - loss: 0.1313 - accuracy: 0.9659 - val_loss: 0.2807 - val_accuracy: 0.8871
Epoch 8/20
30/30 [=====] - 1s 42ms/step - loss: 0.1134 - accuracy: 0.9729 - val_loss: 0.2854 - val_accuracy: 0.8854
Epoch 9/20
30/30 [=====] - 2s 55ms/step - loss: 0.0984 - accuracy: 0.9782 - val_loss: 0.2924 - val_accuracy: 0.8840
Epoch 10/20
30/30 [=====] - 1s 47ms/step - loss: 0.0859 - accuracy: 0.9833 - val_loss: 0.3012 - val_accuracy: 0.8823
Epoch 11/20
30/30 [=====] - 1s 35ms/step - loss: 0.0750 - accuracy: 0.9870 - val_loss: 0.3114 - val_accuracy: 0.8806
Epoch 12/20
30/30 [=====] - 1s 33ms/step - loss: 0.0662 - accuracy: 0.9897 - val_loss: 0.3213 - val_accuracy: 0.8823
Epoch 13/20
30/30 [=====] - 1s 30ms/step - loss: 0.0580 - accuracy: 0.9919 - val_loss: 0.3320 - val_accuracy: 0.8814
Epoch 14/20
30/30 [=====] - 1s 35ms/step - loss: 0.0513 - accuracy: 0.9934 - val_loss: 0.3418 - val_accuracy: 0.8798
Epoch 15/20
30/30 [=====] - 1s 43ms/step - loss: 0.0456 - accuracy: 0.9949 - val_loss: 0.3530 - val_accuracy: 0.8788
Epoch 16/20
30/30 [=====] - 1s 42ms/step - loss: 0.0407 - accuracy: 0.9961 - val_loss: 0.3632 - val_accuracy: 0.8775
Epoch 17/20
30/30 [=====] - 1s 42ms/step - loss: 0.0364 - accuracy: 0.9969 - val_loss: 0.3741 - val_accuracy: 0.8771
Epoch 18/20
30/30 [=====] - 1s 36ms/step - loss: 0.0328 - accuracy: 0.9974 - val_loss: 0.3853 - val_accuracy: 0.8766
Epoch 19/20
30/30 [=====] - 1s 32ms/step - loss: 0.0293 - accuracy: 0.9986 - val_loss: 0.3946 - val_accuracy: 0.8747
Epoch 20/20
30/30 [=====] - 1s 47ms/step - loss: 0.0264 - accuracy: 0.9991 - val_loss: 0.4056 - val_accuracy: 0.8753

```

```

history1_3 = model1_3.fit(partial_x_train,
                          partial_y_train,
                          epochs=20,
                          batch_size=512,
                          validation_data=(x_val, y_val))

```

```

Epoch 1/20
30/30 [=====] - 6s 158ms/step - loss: 0.5474 - accuracy: 0.7514 - val_loss: 0.3953 - val_accuracy: 0.8553
Epoch 2/20
30/30 [=====] - 2s 58ms/step - loss: 0.2852 - accuracy: 0.9023 - val_loss: 0.2877 - val_accuracy: 0.8885
Epoch 3/20
30/30 [=====] - 2s 55ms/step - loss: 0.1842 - accuracy: 0.9375 - val_loss: 0.2795 - val_accuracy: 0.8907
Epoch 4/20
30/30 [=====] - 1s 31ms/step - loss: 0.1300 - accuracy: 0.9589 - val_loss: 0.2990 - val_accuracy: 0.8853
Epoch 5/20
30/30 [=====] - 1s 31ms/step - loss: 0.0928 - accuracy: 0.9728 - val_loss: 0.3279 - val_accuracy: 0.8825
Epoch 6/20

```

```

30/30 [=====] - 1s 32ms/step - loss: 0.0648 - accuracy: 0.9839 - val_loss: 0.3715 - val_accuracy: 0.8806
Epoch 7/20
30/30 [=====] - 1s 33ms/step - loss: 0.0455 - accuracy: 0.9896 - val_loss: 0.4202 - val_accuracy: 0.8763
Epoch 8/20
30/30 [=====] - 2s 51ms/step - loss: 0.0313 - accuracy: 0.9945 - val_loss: 0.4667 - val_accuracy: 0.8740
Epoch 9/20
30/30 [=====] - 1s 41ms/step - loss: 0.0202 - accuracy: 0.9980 - val_loss: 0.5151 - val_accuracy: 0.8716
Epoch 10/20
30/30 [=====] - 2s 53ms/step - loss: 0.0135 - accuracy: 0.9993 - val_loss: 0.5579 - val_accuracy: 0.8698
Epoch 11/20
30/30 [=====] - 1s 40ms/step - loss: 0.0092 - accuracy: 0.9997 - val_loss: 0.5984 - val_accuracy: 0.8676
Epoch 12/20
30/30 [=====] - 1s 51ms/step - loss: 0.0066 - accuracy: 0.9999 - val_loss: 0.6347 - val_accuracy: 0.8682
Epoch 13/20
30/30 [=====] - 2s 65ms/step - loss: 0.0050 - accuracy: 0.9999 - val_loss: 0.6665 - val_accuracy: 0.8678
Epoch 14/20
30/30 [=====] - 1s 37ms/step - loss: 0.0039 - accuracy: 0.9999 - val_loss: 0.6968 - val_accuracy: 0.8663
Epoch 15/20
30/30 [=====] - 1s 51ms/step - loss: 0.0031 - accuracy: 0.9999 - val_loss: 0.7203 - val_accuracy: 0.8676
Epoch 16/20
30/30 [=====] - 1s 40ms/step - loss: 0.0026 - accuracy: 0.9999 - val_loss: 0.7444 - val_accuracy: 0.8673
Epoch 17/20
30/30 [=====] - 1s 33ms/step - loss: 0.0022 - accuracy: 0.9999 - val_loss: 0.7664 - val_accuracy: 0.8659
Epoch 18/20
30/30 [=====] - 1s 32ms/step - loss: 0.0018 - accuracy: 0.9999 - val_loss: 0.7875 - val_accuracy: 0.8655
Epoch 19/20
30/30 [=====] - 1s 35ms/step - loss: 0.0016 - accuracy: 0.9999 - val_loss: 0.8074 - val_accuracy: 0.8649
Epoch 20/20
30/30 [=====] - 1s 41ms/step - loss: 0.0014 - accuracy: 0.9999 - val_loss: 0.8239 - val_accuracy: 0.8649

```

plotting training vs validation loss

```

historyp1_1 = history1_1.history
historyp1_1.keys()

```

```

dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

```

```

historyp1_3 = history1_1.history
historyp1_3.keys()

```

```

dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])

```


```

historyp1_1 = history1_1.history
loss_values1 = historyp1_1["loss"]
val_loss_values1 = historyp1_1["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values1, "bo", label="Training loss")
plt.plot(epochs, val_loss_values1, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()


```

Training and validation loss


```
historyp1_3 = history1_3.history
loss_values3 = historyp1_3["loss"]
val_loss_values3 = historyp1_3["val_loss"]
epochs = range(1, len(loss_values3) + 1)
plt.plot(epochs, loss_values3, "bo", label="Training loss")
plt.plot(epochs, val_loss_values3, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()
```



```
plt.clf()
acc1 = historyp1_1["accuracy"]
val_acc1 = historyp1_1["val_accuracy"]
plt.plot(epochs, acc1, "bo", label="Training acc")
plt.plot(epochs, val_acc1, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```




```
plt.clf()
acc3 = historyp1_3["accuracy"]
val_acc3 = historyp1_3["val_accuracy"]
plt.plot(epochs, acc3, "bo", label="Training acc")
plt.plot(epochs, val_acc3, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()
```



▼ 2 For the hidden layers we are using nodes 32 units, 64 units

```
model2 = keras.Sequential([
    layers.Dense(32, activation="relu"),
    layers.Dense(64, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])
```

```
model2.compile(optimizer="adam",
               loss="binary_crossentropy",
               metrics=["accuracy"])
```

```
hist2 = model2.fit(partial_x_train,
                   partial_y_train,
                   epochs=20,
                   batch_size=512,
                   validation_data=(x_val, y_val))
```

```
Epoch 1/20
30/30 [=====] - 3s 74ms/step - loss: 0.5223 - accuracy: 0.7798 - val_loss: 0.3463 - val_accuracy: 0.8686
Epoch 2/20
30/30 [=====] - 1s 40ms/step - loss: 0.2485 - accuracy: 0.9105 - val_loss: 0.2777 - val_accuracy: 0.8898
Epoch 3/20
30/30 [=====] - 1s 39ms/step - loss: 0.1621 - accuracy: 0.9453 - val_loss: 0.2841 - val_accuracy: 0.8873
Epoch 4/20
30/30 [=====] - 2s 55ms/step - loss: 0.1130 - accuracy: 0.9649 - val_loss: 0.3099 - val_accuracy: 0.8815
Epoch 5/20
30/30 [=====] - 2s 52ms/step - loss: 0.0786 - accuracy: 0.9786 - val_loss: 0.3436 - val_accuracy: 0.8809
Epoch 6/20
30/30 [=====] - 1s 42ms/step - loss: 0.0530 - accuracy: 0.9888 - val_loss: 0.3869 - val_accuracy: 0.8760
Epoch 7/20
30/30 [=====] - 1s 40ms/step - loss: 0.0341 - accuracy: 0.9945 - val_loss: 0.4333 - val_accuracy: 0.8727
Epoch 8/20
30/30 [=====] - 2s 61ms/step - loss: 0.0217 - accuracy: 0.9975 - val_loss: 0.4765 - val_accuracy: 0.8713
Epoch 9/20
30/30 [=====] - 2s 51ms/step - loss: 0.0129 - accuracy: 0.9997 - val_loss: 0.5196 - val_accuracy: 0.8677
Epoch 10/20
30/30 [=====] - 1s 38ms/step - loss: 0.0080 - accuracy: 0.9998 - val_loss: 0.5520 - val_accuracy: 0.8689
Epoch 11/20
30/30 [=====] - 1s 38ms/step - loss: 0.0054 - accuracy: 0.9999 - val_loss: 0.5815 - val_accuracy: 0.8679
```

```

Epoch 12/20
30/30 [=====] - 1s 46ms/step - loss: 0.0038 - accuracy: 0.9999 - val_loss: 0.6093 - val_accuracy: 0.8682
Epoch 13/20
30/30 [=====] - 1s 41ms/step - loss: 0.0027 - accuracy: 0.9999 - val_loss: 0.6334 - val_accuracy: 0.8674
Epoch 14/20
30/30 [=====] - 1s 41ms/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 0.6598 - val_accuracy: 0.8661
Epoch 15/20
30/30 [=====] - 2s 58ms/step - loss: 0.0014 - accuracy: 1.0000 - val_loss: 0.6847 - val_accuracy: 0.8670
Epoch 16/20
30/30 [=====] - 1s 40ms/step - loss: 0.0011 - accuracy: 1.0000 - val_loss: 0.7076 - val_accuracy: 0.8660
Epoch 17/20
30/30 [=====] - 2s 55ms/step - loss: 8.0906e-04 - accuracy: 1.0000 - val_loss: 0.7308 - val_accuracy: 0.8663
Epoch 18/20
30/30 [=====] - 2s 75ms/step - loss: 6.2698e-04 - accuracy: 1.0000 - val_loss: 0.7510 - val_accuracy: 0.8660
Epoch 19/20
30/30 [=====] - 2s 51ms/step - loss: 4.9709e-04 - accuracy: 1.0000 - val_loss: 0.7677 - val_accuracy: 0.8662
Epoch 20/20
30/30 [=====] - 2s 57ms/step - loss: 4.0365e-04 - accuracy: 1.0000 - val_loss: 0.7845 - val_accuracy: 0.8651

```

```

histp2 = hist2.history
loss_values = histp2["loss"]
val_loss_values = histp2["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, "bo", label="Training loss")
plt.plot(epochs, val_loss_values, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()

```

```

plt.clf()
acc = histp2["accuracy"]
val_acc = histp2["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()


```



```

val_loss_values = histp3["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, "bo", label="Training loss")
plt.plot(epochs, val_loss_values, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()


```



```

plt.clf()
acc = histp3["accuracy"]
val_acc = histp3["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()
plt.show()

```



▸ 4 I am using tanh activation instead of relu.

```

model4 = keras.Sequential([
    layers.Dense(16, activation="tanh"),
    layers.Dense(16, activation="tanh"),
    layers.Dense(1, activation="sigmoid")
])

model4.compile(optimizer="adam",
               loss="mse",
               metrics=["accuracy"])


hist4 = model4.fit(partial_x_train,
                  partial_y_train,
                  epochs=20,
                  batch_size=512,
                  validation_data=(x_val, y_val))

Epoch 1/20
30/30 [=====] - 3s 81ms/step - loss: 0.1672 - accuracy: 0.7833 - val_loss: 0.1124 - val_accuracy: 0.8655
Epoch 2/20
30/30 [=====] - 2s 56ms/step - loss: 0.0792 - accuracy: 0.9085 - val_loss: 0.0868 - val_accuracy: 0.8865
Epoch 3/20
30/30 [=====] - 1s 41ms/step - loss: 0.0520 - accuracy: 0.9423 - val_loss: 0.0831 - val_accuracy: 0.8880
Epoch 4/20
30/30 [=====] - 1s 35ms/step - loss: 0.0368 - accuracy: 0.9639 - val_loss: 0.0842 - val_accuracy: 0.8857
Epoch 5/20
30/30 [=====] - 1s 39ms/step - loss: 0.0267 - accuracy: 0.9765 - val_loss: 0.0872 - val_accuracy: 0.8806
Epoch 6/20
30/30 [=====] - 1s 42ms/step - loss: 0.0198 - accuracy: 0.9841 - val_loss: 0.0914 - val_accuracy: 0.8752
Epoch 7/20
30/30 [=====] - 1s 35ms/step - loss: 0.0152 - accuracy: 0.9885 - val_loss: 0.0929 - val_accuracy: 0.8774
Epoch 8/20
30/30 [=====] - 1s 32ms/step - loss: 0.0119 - accuracy: 0.9915 - val_loss: 0.0962 - val_accuracy: 0.8729
Epoch 9/20
30/30 [=====] - 1s 43ms/step - loss: 0.0092 - accuracy: 0.9934 - val_loss: 0.0983 - val_accuracy: 0.8729
Epoch 10/20
30/30 [=====] - 2s 53ms/step - loss: 0.0076 - accuracy: 0.9944 - val_loss: 0.1006 - val_accuracy: 0.8714
Epoch 11/20
30/30 [=====] - 2s 62ms/step - loss: 0.0066 - accuracy: 0.9948 - val_loss: 0.1023 - val_accuracy: 0.8692
Epoch 12/20
30/30 [=====] - 1s 46ms/step - loss: 0.0058 - accuracy: 0.9954 - val_loss: 0.1033 - val_accuracy: 0.8697
Epoch 13/20
30/30 [=====] - 2s 53ms/step - loss: 0.0053 - accuracy: 0.9957 - val_loss: 0.1043 - val_accuracy: 0.8697
Epoch 14/20
30/30 [=====] - 1s 37ms/step - loss: 0.0049 - accuracy: 0.9959 - val_loss: 0.1056 - val_accuracy: 0.8673
Epoch 15/20
30/30 [=====] - 1s 32ms/step - loss: 0.0047 - accuracy: 0.9959 - val_loss: 0.1064 - val_accuracy: 0.8675
Epoch 16/20
30/30 [=====] - 1s 38ms/step - loss: 0.0045 - accuracy: 0.9960 - val_loss: 0.1069 - val_accuracy: 0.8681
Epoch 17/20
30/30 [=====] - 1s 37ms/step - loss: 0.0044 - accuracy: 0.9960 - val_loss: 0.1076 - val_accuracy: 0.8675
Epoch 18/20
30/30 [=====] - 1s 34ms/step - loss: 0.0042 - accuracy: 0.9961 - val_loss: 0.1085 - val_accuracy: 0.8677
Epoch 19/20
30/30 [=====] - 1s 32ms/step - loss: 0.0041 - accuracy: 0.9962 - val_loss: 0.1095 - val_accuracy: 0.8663
Epoch 20/20
30/30 [=====] - 1s 39ms/step - loss: 0.0039 - accuracy: 0.9965 - val_loss: 0.1101 - val_accuracy: 0.8653


histp4 = hist4.history
loss_values = histp4["loss"]
val_loss_values = histp4["val_loss"]
epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, "bo", label="Training loss")
plt.plot(epochs, val_loss_values, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.show()

plt.clf()
acc = histp4["accuracy"]
val_acc = histp4["val_accuracy"]
plt.plot(epochs, acc, "bo", label="Training acc")
plt.plot(epochs, val_acc, "b", label="Validation acc")
plt.title("Training and validation accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.legend()

```



<matplotlib.legend.Legend at 0x7eae20aac910>



▼ 5 In our network I am using Dropout Technique.

Double-click (or enter) to edit

```
#I am using the dropout method with two hidden layers that have the ReLu activation function.
from tensorflow import keras
from tensorflow.keras import layers
model5 = keras.Sequential([
    layers.Dense(16, activation="relu"),
    layers.Dropout(0.5),
    layers.Dense(16, activation="relu"),
    layers.Dense(1, activation="sigmoid")
])

model5.compile(optimizer="adam",
               loss="binary_crossentropy",
               metrics=["accuracy"])

hist5 = model5.fit(partial_x_train,
                  partial_y_train,
                  epochs=20,
```

```
batch_size=512,  
validation_data=(x_val, y_val))
```

```
Epoch 1/20  
30/30 [=====] - 4s 99ms/step - loss: 0.6277 - accuracy: 0.6443 - val_loss: 0.5083 - val_accuracy: 0.8402  
Epoch 2/20  
30/30 [=====] - 1s 33ms/step - loss: 0.4561 - accuracy: 0.8172 - val_loss: 0.3690 - val_accuracy: 0.8745  
Epoch 3/20  
30/30 [=====] - 1s 42ms/step - loss: 0.3427 - accuracy: 0.8771 - val_loss: 0.3068 - val_accuracy: 0.8871  
Epoch 4/20  
30/30 [=====] - 1s 42ms/step - loss: 0.2712 - accuracy: 0.9093 - val_loss: 0.2817 - val_accuracy: 0.8907  
Epoch 5/20  
30/30 [=====] - 1s 43ms/step - loss: 0.2284 - accuracy: 0.9275 - val_loss: 0.2949 - val_accuracy: 0.8853  
Epoch 6/20  
30/30 [=====] - 2s 52ms/step - loss: 0.1910 - accuracy: 0.9398 - val_loss: 0.2935 - val_accuracy: 0.8881  
Epoch 7/20  
30/30 [=====] - 2s 56ms/step - loss: 0.1631 - accuracy: 0.9504 - val_loss: 0.2935 - val_accuracy: 0.8860  
Epoch 8/20  
30/30 [=====] - 2s 58ms/step - loss: 0.1320 - accuracy: 0.9609 - val_loss: 0.3131 - val_accuracy: 0.8849  
Epoch 9/20  
30/30 [=====] - 1s 34ms/step - loss: 0.1160 - accuracy: 0.9673 - val_loss: 0.3475 - val_accuracy: 0.8798  
Epoch 10/20  
30/30 [=====] - 1s 47ms/step - loss: 0.0988 - accuracy: 0.9727 - val_loss: 0.3490 - val_accuracy: 0.8830  
Epoch 11/20  
30/30 [=====] - 1s 34ms/step - loss: 0.0877 - accuracy: 0.9744 - val_loss: 0.3655 - val_accuracy: 0.8815  
Epoch 12/20  
30/30 [=====] - 2s 53ms/step - loss: 0.0740 - accuracy: 0.9801 - val_loss: 0.3938 - val_accuracy: 0.8801  
Epoch 13/20  
30/30 [=====] - 1s 44ms/step - loss: 0.0646 - accuracy: 0.9839 - val_loss: 0.4348 - val_accuracy: 0.8774  
Epoch 14/20  
30/30 [=====] - 1s 42ms/step - loss: 0.0559 - accuracy: 0.9863 - val_loss: 0.4474 - val_accuracy: 0.8800  
Epoch 15/20  
30/30 [=====] - 1s 43ms/step - loss: 0.0493 - accuracy: 0.9872 - val_loss: 0.4780 - val_accuracy: 0.8792  
Epoch 16/20  
30/30 [=====] - 2s 53ms/step - loss: 0.0443 - accuracy: 0.9895 - val_loss: 0.5023 - val_accuracy: 0.8789  
Epoch 17/20  
30/30 [=====] - 1s 50ms/step - loss: 0.0415 - accuracy: 0.9900 - val_loss: 0.5030 - val_accuracy: 0.8802  
Epoch 18/20  
30/30 [=====] - 1s 41ms/step - loss: 0.0374 - accuracy: 0.9907 - val_loss: 0.5100 - val_accuracy: 0.8793  
Epoch 19/20  
30/30 [=====] - 1s 41ms/step - loss: 0.0379 - accuracy: 0.9903 - val_loss: 0.5435 - val_accuracy: 0.8772  
Epoch 20/20  
30/30 [=====] - 1s 41ms/step - loss: 0.0318 - accuracy: 0.9919 - val_loss: 0.5547 - val_accuracy: 0.8794
```

```
#Creating training vs. validation graphs Training vs. validation accuracy and loss
```

```
import matplotlib.pyplot as plt  
histp5 = hist5.history  
loss_values = histp5["loss"]  
val_loss_values = histp5["val_loss"]  
epochs = range(1, len(loss_values) + 1)  
plt.plot(epochs, loss_values, "bo", label="Training loss")  
plt.plot(epochs, val_loss_values, "b", label="Validation loss")  
plt.title("Training and validation loss")  
plt.xlabel("Epochs")  
plt.ylabel("Loss")  
plt.legend()  
plt.show()
```

```
plt.clf()  
acc = histp5["accuracy"]  
val_acc = histp5["val_accuracy"]  
plt.plot(epochs, acc, "bo", label="Training acc")  
plt.plot(epochs, val_acc, "b", label="Validation acc")  
plt.title("Training and validation accuracy")  
plt.xlabel("Epochs")  
plt.ylabel("Accuracy")  
plt.legend()  
plt.show()
```

